

Network Clients

Benjamin Brewster

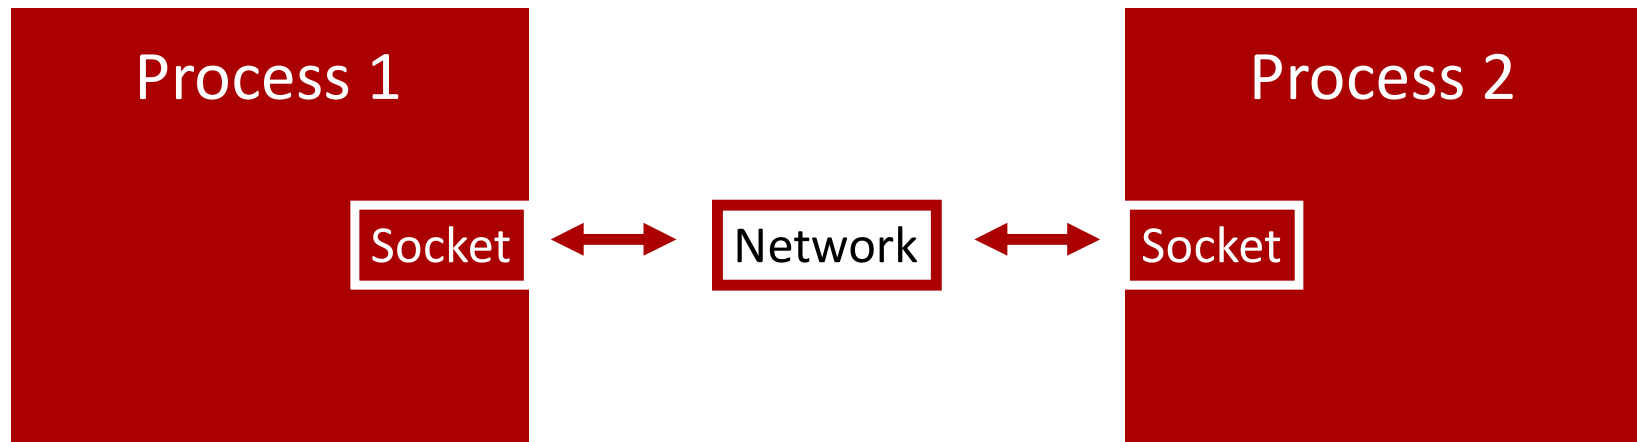
Except as noted, all images copyrighted with Creative Commons licenses,
with attributions given whenever available

Berkeley Sockets

- Developed in the early 1980s for BSD Unix under a grant from DARPA
- The *de facto* network communication method across local area networks (LAN) and the internet.
 - Other transmission methods exist, but they require different transport protocols.
 - i.e., you'd have to write your own version of TCP for a different network protocol
 - Seriously, don't do this. Just use sockets
 - Decades of research! Thousands of scientists, academics, engineers, and hobbyists!
 - Think of the children!

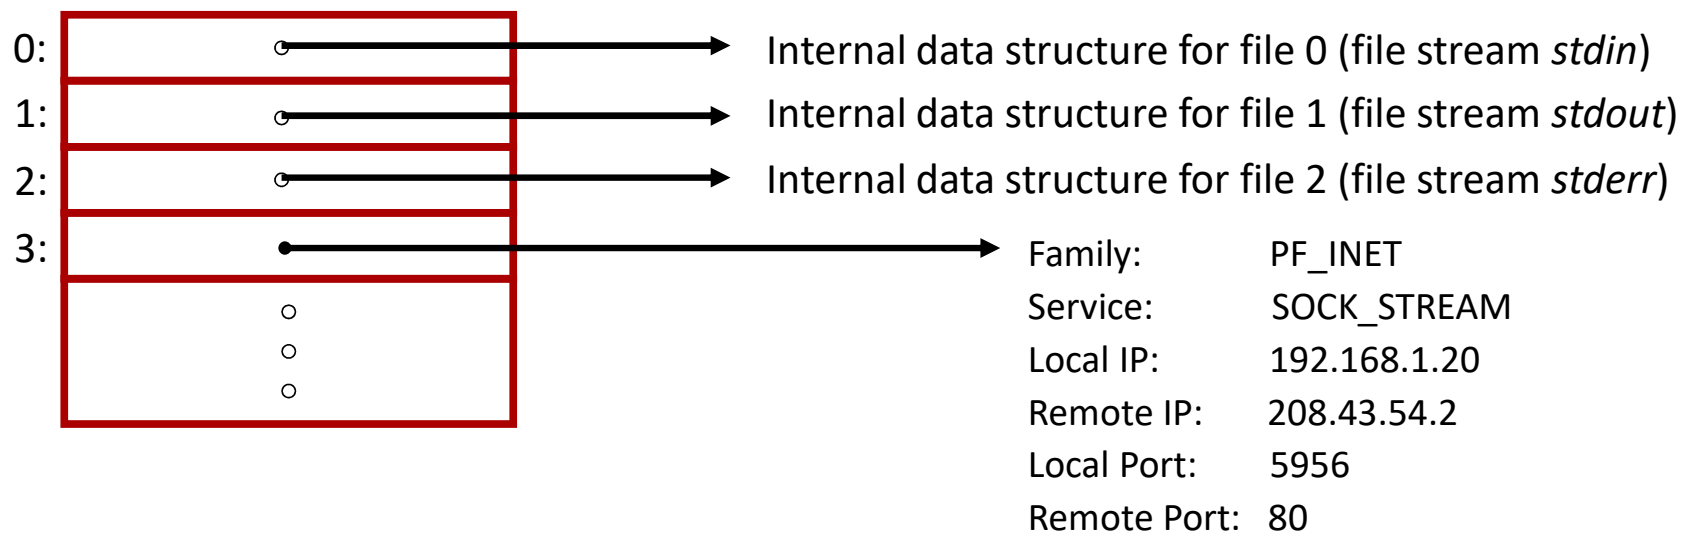


Network Sockets



- Berkeley Socket API
 - A "socket" is the endpoint of a communication link between two processes
 - The socket API treats network connections like *files* as much as possible

File Descriptor Table - Sockets Show Up as Files

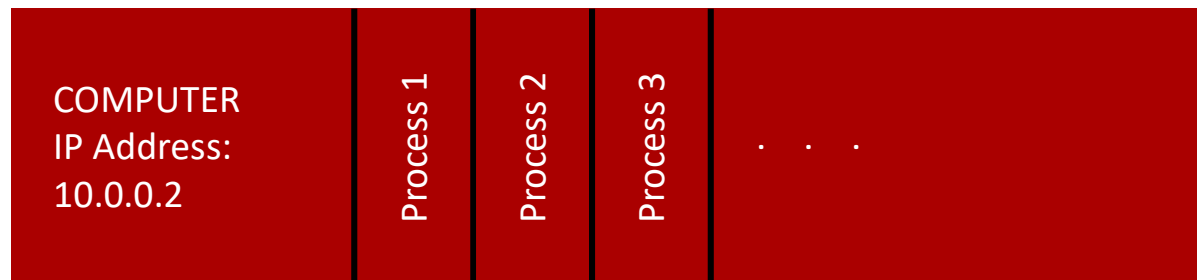


- The file descriptor number returned by `open()` is an index into an array of pointers to internal OS data structures
- Sockets are added to this table of descriptors in the same way



Multiple Process Communication

- Many different processes can be running on one computer
- However, an IP address only identifies the interface on the computer, not the process
- How do we know which process is communicating at that particular interface's IP address?



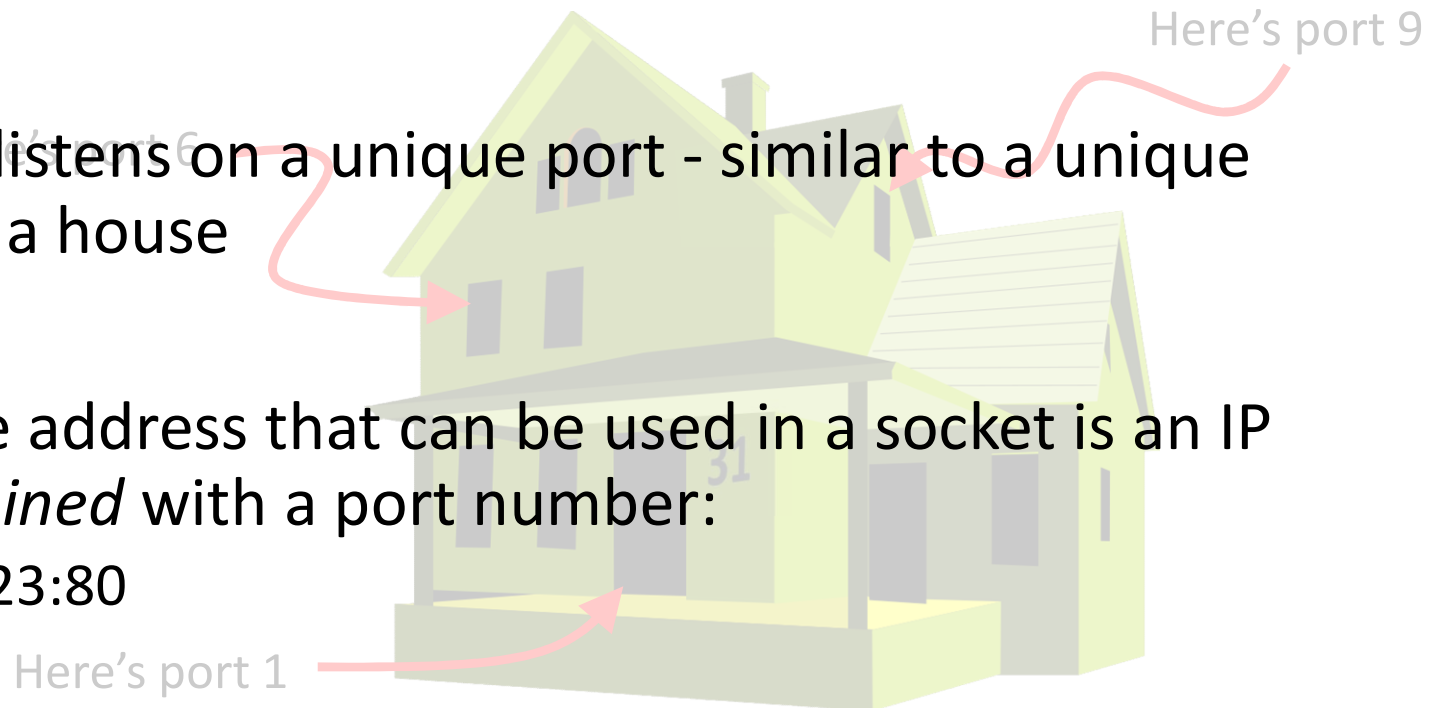
Ports

- This house has address 31



Ports

- *Ports* are used to reach a specific process on a machine
- Each process listens on a unique port - similar to a unique entrance into a house
- So a complete address that can be used in a socket is an IP address *combined* with a port number:
 - 43.144.31.223:80



Socket Documentation

- Most socket related man pages are in the "3n" section
 - `man -s 3n socket`
 - `man -k socket`
- All the info you need to use the network library is scattered across different man pages
- It's definitely best to work from a known good network program starting point! Stay tuned!



Creating and Connecting a Socket on the Client

- Process:
 - 1) Create the socket endpoint with `socket()`
 - 2) Connect the socket to the server with `connect()`
 - 3) Use `read()` and `write()`, or `send()` and `recv()`, to transfer data to and from the socket (which is sent automatically to and from the socket on the server)
- Sockets act like files, in that you can `read()` and `write()` to them
- `send()` and `recv()` are specialized, and can take special flags



Creating the Socket

```
int socket(int domain, int type, int protocol);
```

Returns file
descriptor or -1

For general-purpose sockets that can
connect across a network, use AF_INET
For sockets that are used ONLY for
same-machine IPC, use AF_UNIX

For TCP, use SOCK_STREAM
For UDP, use SOCK_DGRAM

Use 0 for normal behavior

```
int socketFD = socket(AF_INET, SOCK_STREAM, 0);  
if (socketFD == -1) {  
    perror("Hull breach: socket()"); exit(1);  
}
```

Connecting the Socket to an Address

```
int connect(int sockfd, struct sockaddr* address, size_t address_size);
```

Returns 0 on success, -1 on failure

Socket you want to connect

A struct that holds the *address* of where you're connecting, plus other settings;
More on this coming up

The size of the address struct

```
if (connect(socketFD, (struct sockaddr*)&serverAddress, sizeof(serverAddress)))  
{  
    perror("Hull breach: connect()"); exit(1);  
}
```

Filling the Address Struct: IP Address

- Getting the actual address into a form `connect()` can use it is tricky:

```
struct sockaddr_in serverAddress;  
  
serverAddress.sin_family = AF_INET;  
serverAddress.sin_port = htons(7000);  
serverAddress.sin_addr.s_addr = inet_addr("192.168.1.1");
```

`htons()` : **host-to-network-short**

Converts from *host/PC* byte order (LSB) to *network* byte order (MSB)

PCs store bytes with smallest digit first, but networks expect largest digit first

`inet_addr()` converts a standard dotted IP address string into an *integer* format that `sockaddr_in` requires

Filling the Address Struct: Domain Name

- Client connecting to server:

```
struct sockaddr_in serverAddress;  
struct hostent* serverHostInfo;
```

Do a DNS lookup and return address information

```
serverHostInfo = gethostbyname("www.oregonstate.edu");  
if (serverHostInfo == NULL) {  
    fprintf(stderr, "could not resolve server host name\n");  
    exit(1);  
}
```

This will be used to connect to a server on port 80

```
serverAddress.sin_family = AF_INET;  
serverAddress.sin_port = htons(80);
```

Destination, copying to

```
memcpy((char*)&serverAddress.sin_addr.s_addr,  
       (char*)serverHostInfo->h_addr, serverHostInfo->h_length);
```

Preserve the special arrangement of the bytes in these variables by copying the bytes in the given order, regardless of format, structure, or meaning

Source, copying from

Sending Data

Socket file descriptor

Pointer to data that should be sent

```
ssize_t send(int sockfd, void *message, size_t message_size, int flags);
```

Returns number of bytes sent

Number of bytes to send, starting at address in message

Configuration flags

```
char msg[1024];
```

```
...
```

```
r = send(socketFD, msg, 1024, 0);
```

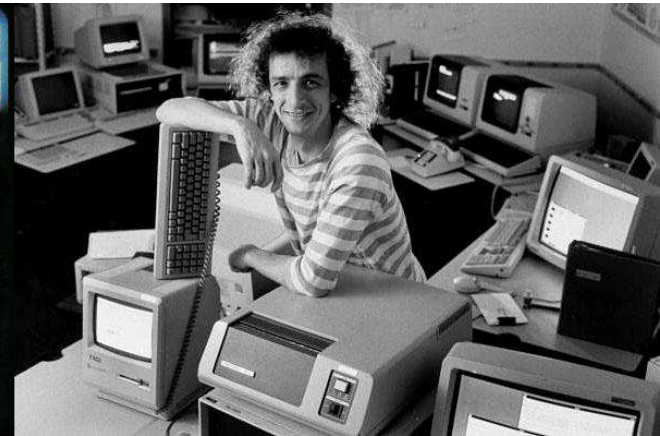
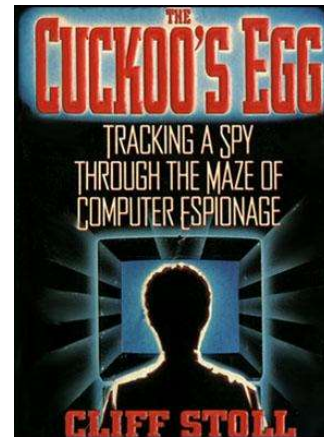
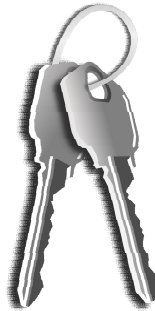
```
if (r < 1024)
```

```
{ } // handle possible error
```

If this happens, you'll have to call `send()` again to send what didn't get sent previously

Sending Data

- Send will block until all the data has been sent, the connection goes away, or a signal handler interrupts the `write()` system call
- Remember that internet connections fail all the time
 - Client intentionally disconnects (STOP button in a web browser)
 - Network failure
 - etc.



Receiving Data

Socket file descriptor

Pointer to where
received data
should be written

```
ssize_t recv(int sockfd, void *buffer, size_t buffer_size, int flags);
```

Returns number
of bytes read

Maximum number of bytes to receive

Configuration flags

```
char buffer[1024];  
memset(buffer, '\0', sizeof(buffer));  
r = recv(socketFD, buffer, sizeof(buffer) - 1, 0);  
if (r < sizeof(buffer) - 1)  
    {} // handle possible error
```

Gray is really, really unlikely

- if $r == -1$, ERROR
- if $0 < r < \text{sizeof}(\text{buffer}) - 1$, there may be more data
- if $r == 0$, sender shut down OR sent a 0-length packet OR 0 bytes were requested

Receiving Data

- Data may arrive in odd size bundles!
- `recv()` or `read()` will return exactly the amount of data that has already arrived
- `recv()` and `read()` will block if the connection is open but *no* data is available
 - So be careful to match what you send with what you receive, or use:

```
fcntl(socketFD, F_SETFL, O_NONBLOCK);
```

...to set the socket to not block if there's no data, but that means you're polling the socket, waiting for data; `select()` would be better (see next lecture!)



Receiving Data - Using Control Codes

- Similar to controlling data being sent through pipes, you can watch for the amount of data coming through `recv()` if you know how much there should be, or use codes:

```
...
char completeMessage[512], readBuffer[10];
memset(completeMessage, '\0', sizeof(completeMessage)); // Clear the buffer


while (strstr(completeMessage, "@@") == NULL) // As long as we haven't found the terminal...
{
    memset(readBuffer, '\0', sizeof(readBuffer)); // Clear the buffer
    r = recv(socketFD, readBuffer, sizeof(readBuffer) - 1, 0); // Get the next chunk
    strcat(completeMessage, readBuffer); // Add that chunk to what we have so far
    printf("PARENT: Message received from child: \"%s\", total: \"%s\"\n", readBuffer, completeMessage);
    if (r == -1) { printf("r == -1\n"); break; } // Check for errors
    if (r == 0) { printf("r == 0\n"); break; }
}

int terminalLocation = strstr(completeMessage, "@@") - completeMessage; // Where is the terminal
completeMessage[terminalLocation] = '\0'; // End the string early to wipe out the terminal
printf("PARENT: Complete string: \"%s\"\n", completeMessage);
...
```

Debugging the Contents of Buffers

- Often, when writing send and receive functions, you'll get garbage. Here's an easy way to actually check what's in a buffer:

```
int x = 0;
printf("CHAR INT\n");
for (x = 0; x < strlen(buffer); x++)
    printf(" %c    %d\n", buffer[x], buffer[x]);
```



Show all chars up to
the first newline

- Or:

```
int x = 0;
printf("CHAR INT\n");
for (x = 0; x < sizeof(buffer); x++)
    printf(" %c    %d\n", buffer[x], buffer[x]);
```



Show all chars in the
entire array

Debugging the Contents of Buffers: Results

CHAR	INT	
o	111	
O	79	
0	48	
	32	Space
l	108	
	10	New line
L	76	
1	49	
	0	Null terminator

- Look up these ints in a good ASCII table, like this one:
<http://www.asciitable.com>



```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
```

This is a basic client program that can send and receive.

It is intended to pair with server.c

client.c 1 of 2

[illegible]

client.c 2 of 2

```
// Set up the socket
socketFD = socket(AF_INET, SOCK_STREAM, 0); // Create the socket
if (socketFD < 0) error("CLIENT: ERROR opening socket");

// Connect to server
if (connect(socketFD, (struct sockaddr*)&serverAddress, sizeof(serverAddress)) < 0) // Connect socket to addy
    error("CLIENT: ERROR connecting");

// Get input message from user
printf("CLIENT: Enter text to send to the server, and then hit enter: ");
memset(buffer, '\0', sizeof(buffer)); // Clear out the buffer array
fgets(buffer, sizeof(buffer) - 1, stdin); // Get input from the user, trunc to buffer - 1 chars, leaving \0
buffer[strcspn(buffer, "\n")] = '\0'; // Remove the trailing \n that fgets adds

// Send message to server
charsWritten = send(socketFD, buffer, strlen(buffer), 0); // Write to the server
if (charsWritten < 0) error("CLIENT: ERROR writing to socket");
if (charsWritten < strlen(buffer)) printf("CLIENT: WARNING: Not all data written to socket!\n");

// Get return message from server
memset(buffer, '\0', sizeof(buffer)); // Clear out the buffer again for reuse
charsRead = recv(socketFD, buffer, sizeof(buffer) - 1, 0); // Read data from the socket, leaving \0 at end
if (charsRead < 0) error("CLIENT: ERROR reading from socket");
printf("CLIENT: I received this from the server: \"%s\"\n", buffer);

close(socketFD); // Close the socket
return 0;
}
```

Client/Server Results

```
$ gcc -o client client.c
```

```
$ gcc -o server server.c
```

```
$ ./server 51717 &
```

```
[1] 21094
```

```
$ ./client localhost 51717
```

```
CLIENT: Enter text to send to the server, and then hit enter: AWESOMESAUCE
```

```
SERVER: I received this from the client: "AWESOMESAUCE"
```

```
CLIENT: I received this from the server: "I am the server, and I got your message"
```

```
[1]+  Done                  ./server 51717
```

```
$
```

