

Network Servers

Benjamin Brewster


Except as noted, all images copyrighted with Creative Commons licenses,
with attributions given whenever available

Network Servers

- This lecture covers
 - Setting up network sockets and connecting clients to them
 - Demo working server code
 - Server concurrency methodologies
 - Knowing when data is available




Non-Concurrent Server Connection Overview




- A server socket listens on a specified port
- Many different clients may be connecting to that port
- The server needs to *differentiate* between and *communicate* with each client

Non-Concurrent Server Connection Overview




- Step 1: The server chooses the next connection to deal with

Non-Concurrent Server Connection Overview




- Step 2: The server creates a new file descriptor for the chosen connection to exclusively use
- Note that the port doesn't change!

Non-Concurrent Server Connection Overview



- Step 3: The chosen client begins communication with the new file descriptor created by the server
- New connections can now be handled again by listenSocketFD

Connection Differentiation



- The OS network layer on the server *differentiates each connection* by using four pieces of the addresses, routing the network packets to the correct socket/FD based on:
 - Server IP, Server Port, Client IP, Client Port

Server Sockets API

- Server Procedure:

1. Create a network socket/FD with `socket()`
2. Bind the socket to a port number with `bind()`
3. Start listening for connections on that socket/port with `listen()`
4. Loop and accept connections on that socket/port with `accept()`, connecting them to *new* sockets for each connection's exclusive use
5. Read and write data to and from the *newly* created sockets for connected and accepted clients using `send()` & `recv()` or `read()` & `write()`



Creating the Socket - Same Method as Client

```
int socket(int domain, int type, int protocol);
```

Returns file
descriptor or -1

For general-purpose sockets that can
connect across a network, use AF_INET
For sockets that are used ONLY for
same-machine IPC, use AF_UNIX

For TCP, use SOCK_STREAM
For UDP, use SOCK_DGRAM

Use 0 for normal behavior

```
int socketFD = socket(AF_INET, SOCK_STREAM, 0);  
if (socketFD < 0) {  
    perror("Hull breach: socket()"); exit(1);  
}
```

Filling the Address Struct for the Server

- Set the address struct so that it accepts connections from any IP address, or just one, and which specific port it will be available on:

```
struct sockaddr_in serverAddress;  
  
serverAddress.sin_family = AF_INET;  
serverAddress.sin_port = htons(80);  
serverAddress.sin_addr.s_addr = INADDR_ANY;
```

`htons()`: **host-to-network-short**

Converts from *host/PC* byte order (LSB) to *network* byte order (MSB)

PCs store bytes with smallest digit first, but networks expect largest digit first

Allows connections from any IP address

Bind the Socket to a Port

- Ports allow multiple processes running on a single machine to communicate across the network from only a single IP address
- A server process has to choose a port where clients can contact it on
- `bind()` associates the chosen port with a socket already created with the `socket()` command
- Subsequent calls to `bind()` using an already-bound socket will fail
- Even after the sockets are all closed, the OS does not immediately release the port; you'll need to wait many seconds for it to be available again for reuse



Binding the Socket

```
int bind(int sockfd, struct sockaddr *address, size_t add_len);
```

Returns 0 on
success or -1
on error

The socket we're
binding to the port

The network address struct,
which identifies which port
this socket will use

The size of the
address struct

```
if (bind(listenSocketFD, (struct sockaddr*)&serverAddress, sizeof(serverAddress)) < 0)
{
    perror("Hull breach: bind()"); exit(1);
}
```

Listening for Connections

- The server will ignore any connection attempts until you tell the socket to start listening for connections with `listen()`
- Once this has been done, the socket will begin queuing up connection requests until it reaches the connection queue limit

```
int listen(int sockfd, int queue_size);
```

Returns 0 on success or -1 on error

The socket we're enabling for connections

Maximum number of connections to queue

```
if (listen(sockfd, 5) < 0) {  
    perror("Hull breach: listen()"); exit(1);  
}
```

Loop and Accept

- Servers generally run continually, waiting for clients to contact them
- Thus a server has an "infinite loop" that continually processes connections from clients
- The `accept()` function takes the next connection off of the listen queue for a socket, or blocks the process until a connection request arrives



Accepting Connections

```
int accept(int sockfd, struct sockaddr* address, size_t &add_len);
```

Returns file descriptor for new connection or -1 on error

The socket we're going to get a connection from

Network address struct, into which connecting *client* information will be written

The size of the address struct

```
int establishedConnectionFD = accept(listenSocketFD,  
                                     (struct sockaddr*)&clientAddress,  
                                     &sizeofClientInfo);  
  
if (establishedConnectionFD < 0) { perror("Hull breach: accept()"); exit(1); }
```

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
```

This is a basic server program
that can send and receive.
It is intended to pair with client.c

server.c 1 of 2

```
void error(const char *msg) { perror(msg); exit(1); } // Error function used for reporting issues

int main(int argc, char *argv[])
{
    int listenSocketFD, establishedConnectionFD, portNumber, charsRead;
    socklen_t sizeofClientInfo;
    char buffer[256];
    struct sockaddr_in serverAddress, clientAddress;

    if (argc < 2) { fprintf(stderr, "USAGE: %s port\n", argv[0]); exit(1); } // Check usage & args

    // Set up the address struct for this process (the server)
    memset((char *)&serverAddress, '\0', sizeof(serverAddress)); // Clear out the address struct
    portNumber = atoi(argv[1]); // Get the port number, convert to an integer from a string
    serverAddress.sin_family = AF_INET; // Create a network-capable socket
    serverAddress.sin_port = htons(portNumber); // Store the port number
    serverAddress.sin_addr.s_addr = INADDR_ANY; // Any address is allowed for connection to this process

    // Set up the socket
    listenSocketFD = socket(AF_INET, SOCK_STREAM, 0); // Create the socket
    if (listenSocketFD < 0) error("ERROR opening socket");
```


server.c 2 of 2

```
// Enable the socket to begin listening
if (bind(listenSocketFD, (struct sockaddr *)&serverAddress, sizeof(serverAddress)) < 0) // Connect socket to port
    error("ERROR on binding");
listen(listenSocketFD, 5); // Flip the socket on - it can now receive up to 5 connections

// Accept a connection, blocking if one is not available until one connects
sizeofClientInfo = sizeof(clientAddress); // Get the size of the address for the client that will connect
establishedConnectionFD = accept(listenSocketFD, (struct sockaddr *)&clientAddress, &sizeofClientInfo); // Accept
if (establishedConnectionFD < 0) error("ERROR on accept");

// Get the message from the client and display it
memset(buffer, '\0', 256);
charsRead = recv(establishedConnectionFD, buffer, 255, 0); // Read the client's message from the socket
if (charsRead < 0) error("ERROR reading from socket");
printf("SERVER: I received this from the client: \"%s\"\n", buffer);

// Send a Success message back to the client
charsRead = send(establishedConnectionFD, "I am the server, and I got your message", 39, 0); // Send success back
if (charsRead < 0) error("ERROR writing to socket");
close(establishedConnectionFD); // Close the existing socket which is connected to the client

close(listenSocketFD); // Close the listening socket
return 0;
}
```

Client/Server Results

```
$ gcc -o client client.c
```

```
$ gcc -o server server.c
```

```
$ ./server 51717 &
```

```
[1] 21094
```

```
$ ./client localhost 51717
```

```
CLIENT: Enter text to send to the server, and then hit enter: AWESOMESAUCE
```

```
SERVER: I received this from the client: "AWESOMESAUCE"
```

```
CLIENT: I received this from the server: "I am the server, and I got your message"
```

```
[1]+  Done                  ./server 51717
```

```
$
```



```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
```

This server lives forever by wrapping
the `accept()` section in a while loop
It is also intended to pair with `client.c`

multiserver.c 1 of 2

```
void error(const char *msg) { perror(msg); exit(1); } // Error function used for reporting issues

int main(int argc, char *argv[])
{
    int listenSocketFD, establishedConnectionFD, portNumber, charsRead;
    socklen_t sizeofClientInfo;
    char buffer[256];
    struct sockaddr_in serverAddress, clientAddress;

    if (argc < 2) { fprintf(stderr, "USAGE: %s port\n", argv[0]); exit(1); } // Check usage & args

    // Set up the address struct for this process (the server)
    memset((char *)&serverAddress, '\0', sizeof(serverAddress)); // Clear out the address struct
    portNumber = atoi(argv[1]); // Get the port number, convert to an integer from a string
    serverAddress.sin_family = AF_INET; // Create a network-capable socket
    serverAddress.sin_port = htons(portNumber); // Store the port number
    serverAddress.sin_addr.s_addr = INADDR_ANY; // Any address is allowed for connection to this process

    // Set up the socket
    listenSocketFD = socket(AF_INET, SOCK_STREAM, 0); // Create the socket
    if (listenSocketFD < 0) error("ERROR opening socket");
```

multiserver.c 2 of 2

```
// Enable the socket to begin listening
if (bind(listenSocketFD, (struct sockaddr *)&serverAddress, sizeof(serverAddress)) < 0) // Connect socket to port
    error("ERROR on binding");
listen(listenSocketFD, 5); // Flip the socket on - it can now receive up to 5 connections
while (1) {
    // Accept a connection, blocking if one is not available until one connects
    sizeofClientInfo = sizeof(clientAddress); // Get the size of the address for the client that will connect
    establishedConnectionFD = accept(listenSocketFD, (struct sockaddr *)&clientAddress, &sizeofClientInfo); // Accept
    if (establishedConnectionFD < 0) error("ERROR on accept");
    printf("SERVER: Connected Client at port %d\n", ntohs(clientAddress.sin_port));
    // Get the message from the client and display it
    memset(buffer, '\0', 256);
    charsRead = recv(establishedConnectionFD, buffer, 255, 0); // Read the client's message from the socket
    if (charsRead < 0) error("ERROR reading from socket");
    printf("SERVER: I received this from the client: \"%s\"\n", buffer);

    // Send a Success message back to the client
    charsRead = send(establishedConnectionFD, "I am the server, and I got your message", 39, 0); // Send success back
    if (charsRead < 0) error("ERROR writing to socket");
    close(establishedConnectionFD); // Close the existing socket which is connected to the client
}
close(listenSocketFD); // Close the listening socket
return 0;
}
```

Client/Multiserver Results

Note that the order of the client and server sending text to the terminal is hard to control!

```
$ multiserver 55556 &
```

```
[1] 26889
```

```
$ client localhost 55556
```

```
CLIENT: Enter text to send to the server, and then hit enter: SERVER: Connected Client at port 38422  
My Test!
```

```
SERVER: I received this from the client: "My Test!"
```

```
CLIENT: I received this from the server: "I am the server, and I got your message"
```

```
$ client localhost 55556
```

```
CLIENT: Enter text to send to the server, and then hit enter: SERVER: Connected Client at port 38424  
So much text!!
```

```
SERVER: I received this from the client: "So much text!!"
```

```
CLIENT: I received this from the server: "I am the server, and I got your message"
```

```
$ kill -TERM 26889
```

```
[1]+  Terminated
```

```
multiserver 55556
```




Doing it All at Once

- Many clients may need to both connect and perform tasks at the same time
- We want to minimize:
 - Response time
 - Complexity
- Want to maximize:
 - Throughput (connections serviced / second)
 - Hardware utilization (%CPU usage)
- These are all tradeoffs of each other!
- Let's look at two methods...



Iterative Servers

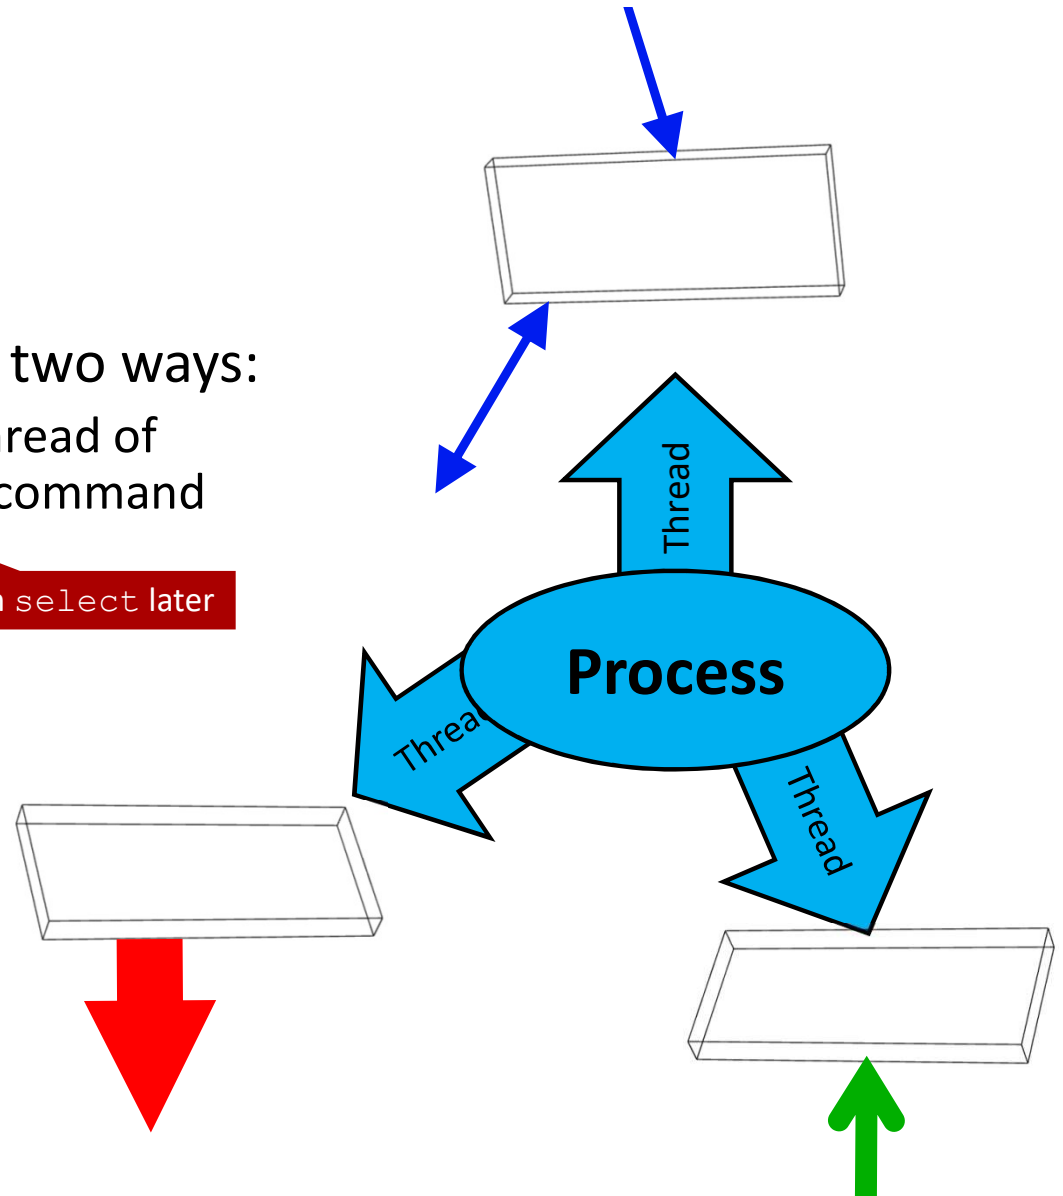
- Iterative
 - Handles only one client at a time
 - Non-preemptive: additional client must wait for all previous requests to complete
- Easy to design, implement, and maintain
- Best when:
 - Request processing time is short
 - No I/O is needed by server
 - Order matters



Concurrent Servers

- Concurrency can be provided in two ways:
 - Apparent concurrency: a single thread of execution, using the `select()` command and non-blocking I/O
 - Real concurrency: multiple threads of execution, or multiple processes, each with one thread

More on `select` later



Apparent Concurrency: Details

- Only one thread, no preemption, using non-blocking I/O
- Whenever an I/O request would block, switch to another connection
- Up to a certain number of clients being server:
 - Maximizes CPU utilization
 - Increases throughput
- Complexity involves tracking connections, choosing the next to run, detecting blocking calls, etc.
- Works well if requests are short



Real Concurrency: Details

- Preemptive
 - Clients can connect anytime to the server, which uses multiple threads or processes to service connections
- Up to a certain number of connections:
 - Maximizes CPU utilization
 - Maximizes response time
 - Increases throughput
- Harder to design, implement, and *maintain*:
 - After too many concurrent connections:
 - Everything gets worse -> server eventually hangs
 - Need to put limits on concurrent connections



More Real Concurrency

- Four different methods
 - Create one process per client connection
 - Create a pool of available processes before clients connect
 - Use only one process, but create one thread per client connection
 - Use only one process, but create a pool of available threads before clients connect



Fork Solution #1



- One process per client connection
- Fork a new process to handle every connection
- Advantages:
 - Simple: minimal shared state to worry about
- Disadvantages:
 - Process creation via `fork()` is slow
 - Context-switching between processes is also slow, but minor compared to `fork()`


Fork Solution #2



- Create a pool of available processes for clients to use
- Advantages:
 - No longer have to fork
 - Have rapid response as long as there is an idle process available
 - Can set the pool size, so that you don't overload the hardware
- Disadvantages:
 - Still have process context switching
 - Managing the pool of processes can be complex


Threads Solutions 1 & 2

- Threads allow multiple concurrent execution contexts within a single process
- Can implement a server as a single process with multiple threads
 - Either one thread per connection, or a pool of threads
- Advantages:
 - Trades process context switches (slow) for thread switches (fast)
 - Shared address space, shared code, shared data, etc.
- Disadvantages:
 - Code must be thread-safe
 - Must always worry about inadvertent data-sharing



select ()

- `select ()` is designed for server-like applications that have many communication channels open at once (like a pool of threads)
 - Data or space may become available at any time on *any* of the channels
 - You want to minimize the delay between when data/space becomes available and your process takes action
- Overview: you call `select ()` with a list of read and/or write file descriptors, and it *returns* when any one of those descriptors becomes readable or writable



select ()

```
int select(  
    int nfd,           // Highest numbered FD + 1  
    fd_set* readfds,   // Input FDs of interest  
    fd_set* writefds,  // Output FDs of interest  
    fd_set* errorfds,  // FDs where exception has occurred  
    struct timeval* timeout // when to time out if nothing happens  
)
```

- The three parameters `readfds`, `writefds`, and `errorfds` are bit masks
 - Each bit of the number refers to one file descriptor
 - Bit 0 is file descriptor 0, bit 1 is file descriptor 1, etc.
- UNIX provides you with macros to manipulate bit masks:
 - `FD_ZERO()` :: Set all bits to 0
 - `FD_SET()` :: Set one specific bit to 1
 - `FD_ISSET()` :: Determine if a specific bit is set to 1
 - `FD_CLR()` :: Set one specific bit to 0

select () Return Values

- -1 if error
- 0 if time out: nothing ready
- Else, the return value is the *number* of file descriptors ready for reading, writing, or have had errors occur

selectDemo.c

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/time.h>
#include <sys/types.h>
#include <unistd.h>

int main(void)
{
    fd_set readFDs;
    struct timeval timeToWait;
    int retval;

    // Watch stdin (FD 0) to see when it has input
    FD_ZERO(&readFDs); // Zero out the set of possible read file descriptors
    FD_SET(0, &readFDs); // Mark only FD 0 as the one we want to pay attention to

    // Wait up to 50 seconds
    timeToWait.tv_sec = 50;
    timeToWait.tv_usec = 0;

    retval = select(1, &readFDs, NULL, NULL, &timeToWait); // Check to see whether any read FDs have data!
                                                         // After select returns, timeToWait is undefined
    if (retval == -1)
        perror("select()");
    else if (retval)
        printf("Data is available now!\n"); // FD_ISSET(0, &readFDs) will return true
    else
        printf("No data within 50 seconds\n");

    return(0);
}
```

This example comes from the
`select()` man page

Together with returning an int, `select()` *also* overwrites your bit masks to show you *which* bits are interesting; you'll have to iterate through them to see which ones are set, though

selectDemo.c

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/time.h>
#include <sys/types.h>
#include <unistd.h>
```

```
int main(void)
{
```

```
    fd_set read_fds;
    struct timeval tv;
    int retval;
```

```
    // Watch for data on
    FD_ZERO(&read_fds);
    FD_SET(0, &read_fds);
```

```
    // Wait until we have
    timeval tv;
    timeToWait(&tv);
```

```
    retval = select(1, &read_fds, NULL, NULL, &tv);
```

```
    if (retval < 0) {
```

```
        // Error
```

```
        return -1;
```


```
    return(0);
```

```
}
```

In marking up this code and testing it, the double slash comments `//` were confusing the copy/paste functions in vim. So, I changed them all to single slashes, pasted it all into vim, and then entered this command:

```
:%s/\//\//g
```

... and it worked, first try.



ad FDs have data!
oWait is undefined

e

selectDemo.c Results

```
$ mkfifo myfifo  
$ selectDemo < myfifo &  
[1] 17013  
$ echo "text" > myfifo  
Data is available now!  
[1]+  Done
```

```
selectDemo < myfifo
```

This hooks the output of the echo command to the input of the selectDemo program through a named pipe!

As soon as both ends are opened, and data is written, the pipe transfers the data

